

```
In [3]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import string
import nltk
from nltk.corpus import stopwords
from wordcloud import WordCloud
nltk.download('stopwords')
import tensorflow as tf
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from sklearn.model_selection import train_test_split
from keras.callbacks import EarlyStopping, ReduceLROnPlateau
import warnings
warnings.filterwarnings('ignore')
```

[nltk\_data] Downloading package stopwords to  
[nltk\_data] C:\Users\Acer\AppData\Roaming\nltk\_data...\nltk\_data]\nUnzipping corpora\stopwords.zip.


```
In [6]: data = pd.read_csv('spam.csv')
data.head(10)
```

	Unnamed: 0	label	text	label_num
0	605	ham	Subject: enron methanol: meter #: 988291Vn...	0
1	2349	ham	Subject: hpl nom for january 9, 2001Vn(see...	0
2	3624	ham	Subject: neon retreat/vrho ho ho , we're ar...	0
3	4685	spam	Subject: photoshop , windows , office , cheap ...	1
4	2030	ham	Subject: re : indian springs/vnthis deal is t...	0
5	2949	ham	Subject: ehonline web address change/vnthis ...	0
6	2793	ham	Subject: spring savings certificate - take 30 ...	0
7	4185	spam	Subject: looking for medication ? we `re the ...	1
8	2641	ham	Subject: noms / actual flow for 2 / 26Vnwe a...	0
9	1870	ham	Subject: nominations for oct . 21 - 23 , 2000V...	0

```
In [7]: data.shape
(5171, 4)
```

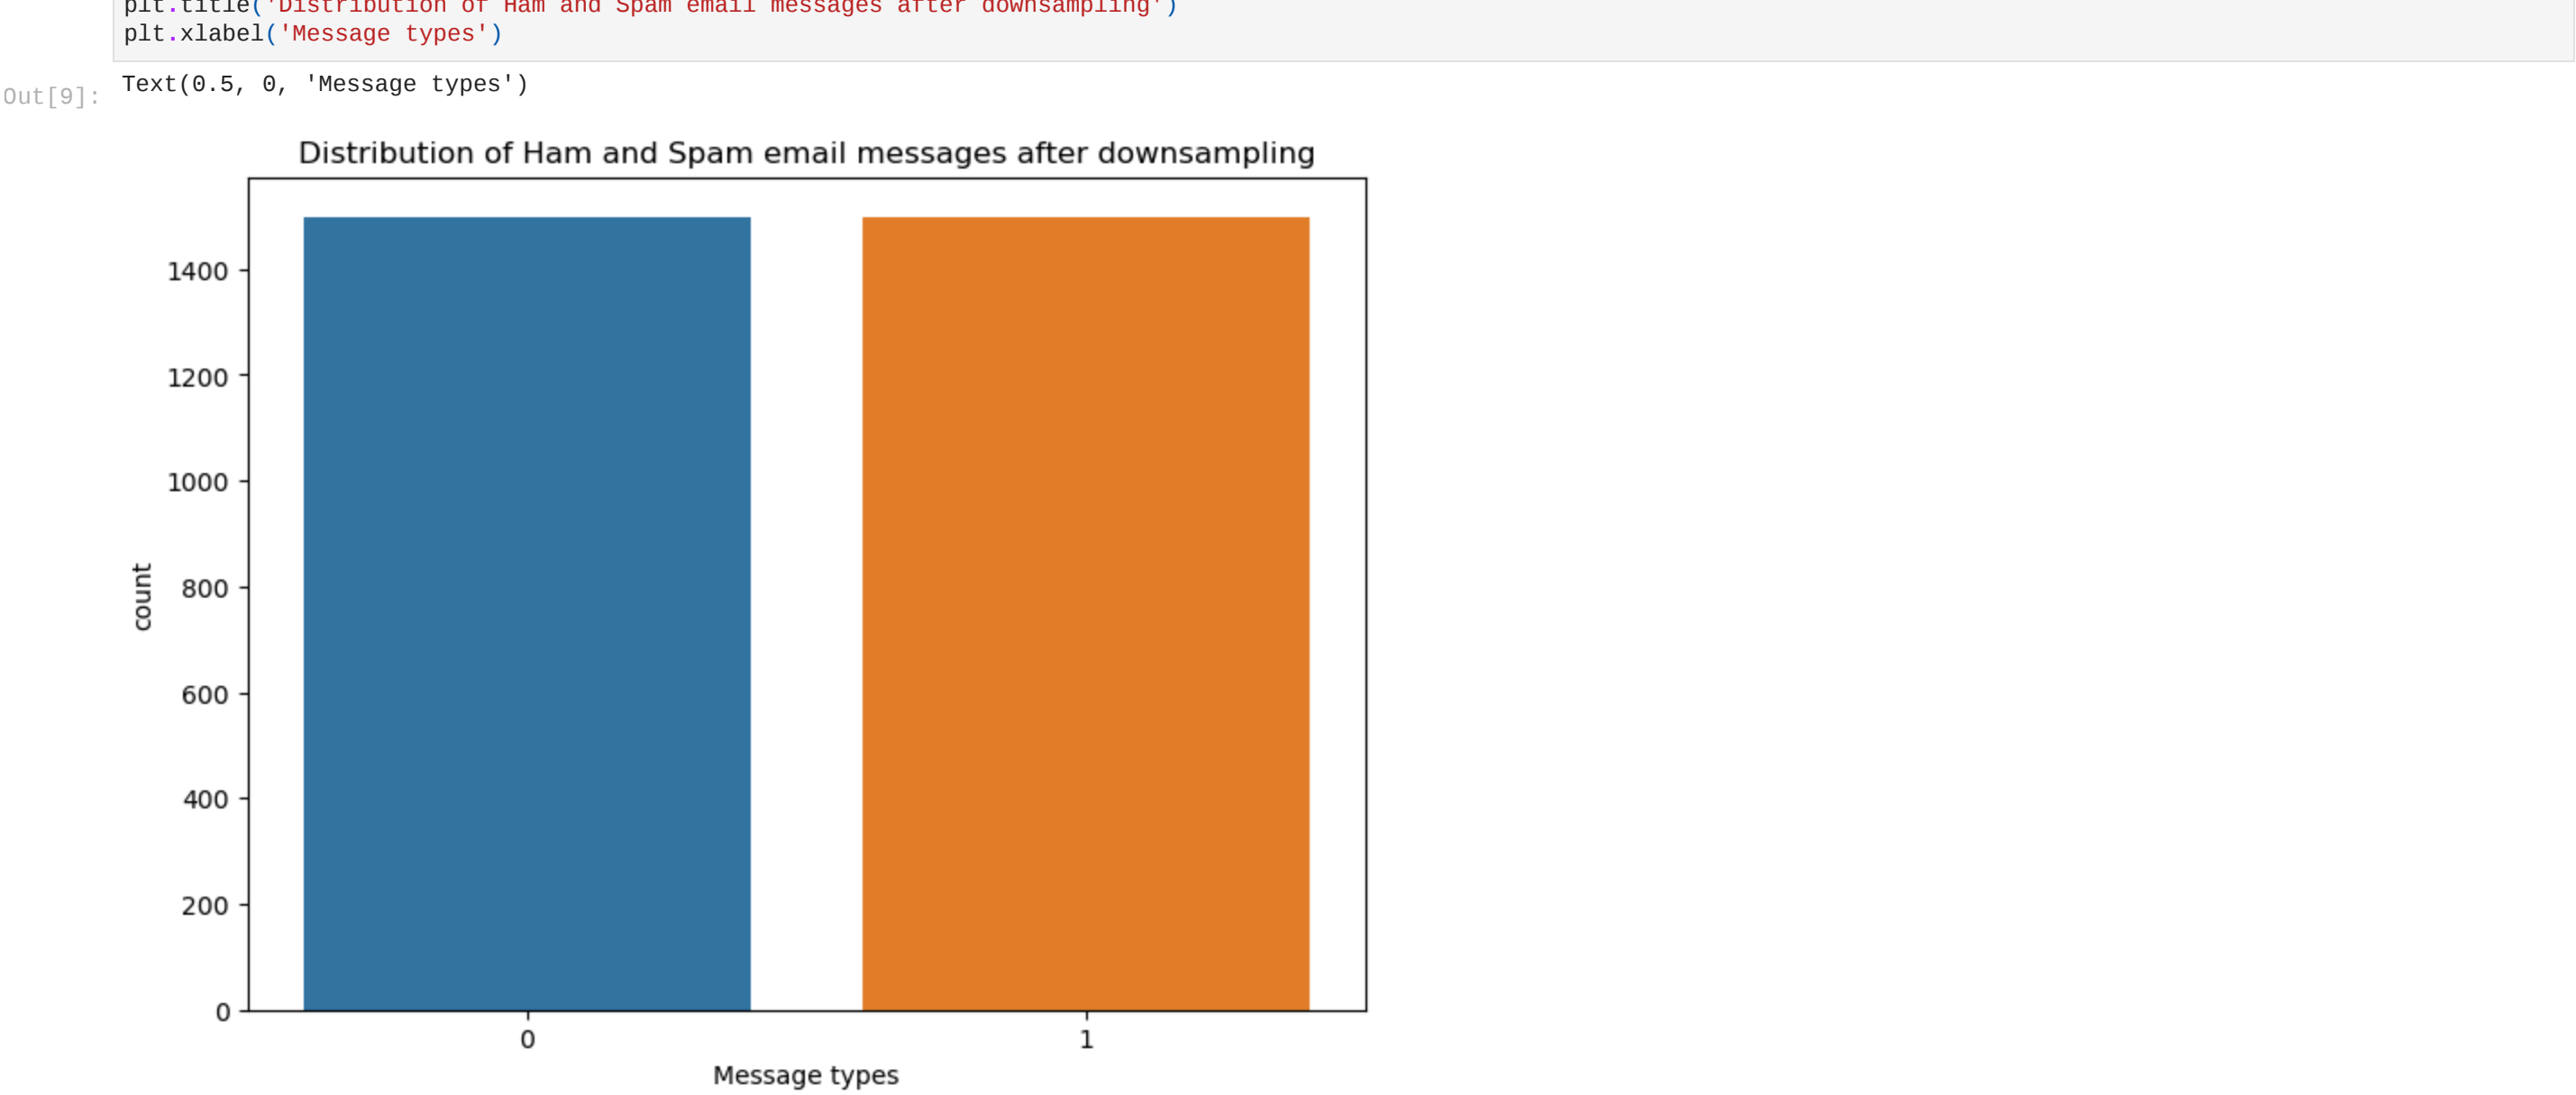
```
Out[7]:
```

```
In [8]: sns.countplot(x='label_num', data=data)
plt.show()
```



```
In [9]: # Downsampling to balance the dataset
ham_msg = data[data.label_num == 0]
spam_msg = data[data.label_num == 1]
ham_msg = ham_msg.sample(n=len(spam_msg), random_state=42)

# Plotting the counts of down sampled dataset
balanced_data = ham_msg.append(spam_msg)\
.reset_index(drop=True)
plt.figure(figsize=(8, 6))
sns.countplot(data = balanced_data, x='label_num')
plt.title('Distribution of Ham and Spam email messages after downsampling')
plt.xlabel('Message types')
Text(0.5, 0, 'Message types')
```



```
In [10]: balanced_data['text'] = balanced_data['text'].str.replace('Subject', '')
balanced_data.head()
```

	Unnamed: 0	label	text	label_num
0	3444	ham	: conoco - big cowboy/vndarren :vni 'm not...	0
1	2982	ham	: feb 01 prod : sale to teco gas processing/vr...	0
2	2711	ham	: california energy crisis/vncalifornia l , s...	0
3	3116	ham	: re : nom / actual volume for april 23 rd/vnwe...	0
4	1314	ham	: easttrans nomination changes effective 8 / 2 ...	0

```
In [11]: punctuations_list = string.punctuation
def remove_punctuations(text):
    temp = str.maketrans('', '', punctuations_list)
    return text.translate(temp)

balanced_data['text'] = balanced_data['text'].apply(lambda x: remove_punctuations(x))
balanced_data.head()
```

	Unnamed: 0	label	text	label_num
0	3444	ham	conoco big cowboy/vndarren vni m not sur...	0
1	2982	ham	feb 01 prod sale to teco gas processing/vrns...	0
2	2711	ham	california energy crisis/vncalifornia l s p...	0
3	3116	ham	re nom actual volume for april 23 rd/vnwe ...	0
4	1314	ham	eastrans nomination changes effective 8 2 0...	0

```
In [12]: def remove_stopwords(text):
stop_words = stopwords.words('english')

imp_words = []

# Storing the important words
for word in str(text).split():
    word = word.lower()

    if word not in stop_words:
        imp_words.append(word)

output = " ".join(imp_words)

return output

balanced_data['text'] = balanced_data['text'].apply(lambda text: remove_stopwords(text))
balanced_data.head()
```

	Unnamed: 0	label	text	label_num
0	3444	ham	conoco big cowboy darren sure help know else a...	0
1	2982	ham	feb 01 prod sale teco gas processing sale deal...	0
2	2711	ham	california energy crisis california l power cr...	0
3	3116	ham	nom actual volume april 23 rd agree eileen pon...	0
4	1314	ham	eastrans nomination changes effective 8 2 00 p...	0


```
In [13]: def plot_word_cloud(data, typ):
email_corpus = " ".join(data['text'])

plt.figure(figsize=(7, 7))

wc = WordCloud(background_color='black',
                max_words=100,
                width=800,
                height=400,
                collocations=False).generate(email_corpus)

plt.imshow(wc, interpolation='bilinear')
plt.title(f'WordCloud for {typ} emails', fontsize=15)
plt.axis('off')
plt.show()

plot_word_cloud(balanced_data[balanced_data['label_num'] == 0], typ='Non-Spam')
plot_word_cloud(balanced_data[balanced_data['label_num'] == 1], typ='Spam')
```



```
In [14]: #train test split
train_X, test_X, train_Y, test_Y = train_test_split(balanced_data['text'],
                                                    balanced_data['label_num'],
                                                    test_size = 0.2,
                                                    random_state = 42)
```

```
In [15]: # Tokenize the text data
tokenizer = Tokenizer()
tokenizer.fit_on_texts(train_X)

# Convert text to sequences
train_sequences = tokenizer.texts_to_sequences(train_X)
test_sequences = tokenizer.texts_to_sequences(test_X)

# Pad sequences to have the same length
max_len = 100 # maximum sequence length
train_sequences = pad_sequences(train_sequences,
                                maxlen=max_len,
                                padding='post',
                                truncating='post')

test_sequences = pad_sequences(test_sequences,
                                maxlen=max_len,
                                padding='post',
                                truncating='post')
```

```
In [16]: # Build the model
model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Embedding(input_dim=len(tokenizer.word_index) + 1,
                                     output_dim=32,
                                     input_length=max_len))

model.add(tf.keras.layers.LSTM(16))
model.add(tf.keras.layers.Dense(32, activation='relu'))
model.add(tf.keras.layers.Dense(1, activation='sigmoid'))

# Print the model summary
model.summary()
```

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 100, 32)	1274912
lstm (LSTM)	(None, 16)	3136
dense (Dense)	(None, 32)	544
dense_1 (Dense)	(None, 1)	33

Total params: 1278625 (4.88 MB)  
Trainable params: 1278625 (4.88 MB)  
Non-trainable params: 0 (0.00 Byte)

```
In [17]: model.compile(loss = tf.keras.losses.BinaryCrossentropy(from_logits = True),
                    metrics = ['accuracy'],
                    optimizer = 'adam')
```

```
In [18]: es = EarlyStopping(patience=3,
                        monitor = 'val_accuracy',
                        restore_best_weights = True)

lr = ReduceLROnPlateau(patience = 2,
                        monitor = 'val_loss',
                        factor = 0.5,
                        verbose = 0)
```

```
In [19]: # Train the model
history = model.fit(train_sequences, train_Y,
                    validation_data=(test_sequences, test_Y),
                    epochs=20,
                    batch_size=32,
                    callbacks = [lr, es])
```

Epoch 1/20	75/75 [=====]	- 8s 70ms/step - loss: 0.6339 - accuracy: 0.6493 - val_loss: 0.3189 - val_accuracy: 0.9600 - lr: 0.0010
Epoch 2/20	75/75 [=====]	- 4s 59ms/step - loss: 0.2170 - accuracy: 0.9462 - val_loss: 0.1568 - val_accuracy: 0.9667 - lr: 0.0010
Epoch 3/20	75/75 [=====]	- 4s 60ms/step - loss: 0.1411 - accuracy: 0.9683 - val_loss: 0.1547 - val_accuracy: 0.9633 - lr: 0.0010
Epoch 4/20	75/75 [=====]	- 4s 59ms/step - loss: 0.1233 - accuracy: 0.9725 - val_loss: 0.1622 - val_accuracy: 0.9617 - lr: 0.0010
Epoch 5/20	75/75 [=====]	- 4s 48ms/step - loss: 0.1307 - accuracy: 0.9691 - val_loss: 0.1402 - val_accuracy: 0.9683 - lr: 0.0010
Epoch 6/20	75/75 [=====]	- 4s 51ms/step - loss: 0.0996 - accuracy: 0.9779 - val_loss: 0.1146 - val_accuracy: 0.9750 - lr: 0.0010
Epoch 7/20	75/75 [=====]	- 4s 51ms/step - loss: 0.1155 - accuracy: 0.9729 - val_loss: 0.1412 - val_accuracy: 0.9617 - lr: 0.0010
Epoch 8/20	75/75 [=====]	- 4s 49ms/step - loss: 0.0859 - accuracy: 0.9821 - val_loss: 0.4893 - val_accuracy: 0.8850 - lr: 0.0010
Epoch 9/20	75/75 [=====]	- 3s 47ms/step - loss: 0.1536 - accuracy: 0.9629 - val_loss: 0.1350 - val_accuracy: 0.9683 - lr: 5.0000e-04

```
In [20]: # Train the model
history = model.fit(train_sequences, train_Y,
                    validation_data=(test_sequences, test_Y),
                    epochs=20,
                    batch_size=32,
                    callbacks = [lr, es])
```

Epoch 1/20	75/75 [=====]	- 3s 39ms/step - loss: 0.0839 - accuracy: 0.9629 - val_loss: 0.0967 - val_accuracy: 0.9800 - lr: 5.0000e-04
Epoch 2/20	75/75 [=====]	- 4s 48ms/step - loss: 0.0654 - accuracy: 0.9595 - val_loss: 0.1275 - val_accuracy: 0.9700 - lr: 5.0000e-04
Epoch 3/20	75/75 [=====]	- 3s 46ms/step - loss: 0.1199 - accuracy: 0.9708 - val_loss: 0.1169 - val_accuracy: 0.9633 - lr: 5.0000e-04
Epoch 4/20	75/75 [=====]	- 3s 46ms/step - loss: 0.1094 - accuracy: 0.9733 - val_loss: 0.1347 - val_accuracy: 0.9617 - lr: 2.5000e-04

```
In [21]: # Evaluate the model
test_loss, test_accuracy = model.evaluate(test_sequences, test_Y)
print('Test Loss : ', test_loss)
print('Test Accuracy : ', test_accuracy)
```

Test Loss : 0.09668896347284317  
Test Accuracy : 0.9800906190734863

```
In [22]: plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



```
In [ ]:
```